

CS180 Homework #6

1. Exercise 19 on page 329

Algorithm

- Define a string as an interleaving only if it is made up of strings x and y where the last character of the string is from either x or y .
- Let $O(i, j)$ be the optimal solution for a string of length $i+j$, where i = last index of our proposed string x' and j = last index of our proposed string y' .
- Let $O(0, 0)$ be a valid interleaved string.
- For string length $L = 1, 2, \dots, n$
 - For all pairs of (i, j) such that $i+j = k$:
 - If $O(i-1, j)$ is a valid string and the last character is $x'[i]$:
 - $O(i, j)$ is a valid string
 - Else if $O(i, j-1)$ is a valid string and the last character is $y'[j]$
 - $O(i, j)$ is a valid string
 - Otherwise, $O(i, j)$ isn't a valid string
- If there is a pair (i, j) where $i+j = n$ and $O(i, j)$ is a valid string, then there exists a valid interleaving.

Proof of Correctness

- Base case:
 - String length = 0 signifies an empty string, which is a valid interleaving
- Inductive hypothesis:
 - Suppose we know the validity for strings of length 1 to $n-1$, given different prefixes of x and y – x' and y' .
 - To find the validity of a string with length n , we consider the last character, which could either come from x' or y' .
 - If the last character came from x' , then we check whether the string of length $n-1$ is a valid interleaving of x' without its last character and y' – if so, then our current string is valid.
 - If the last character came from y' , check whether the string $n-1$ is a valid interleaving of y' without its last character and x' – if so, our current string is valid.
 - Therefore, our algorithm works to find the validity of a string of length n .

Time Complexity

- We are considering all possible string lengths, which is $O(N)$.
- The inner loop for considering all pairs of (i, j) for a set string length also takes $O(N)$. Therefore, the 2 loops combined takes $O(N^2)$ time.
- Total runtime = $O(N^2)$.

2. Exercise 22 on page 330

Algorithm

- Create a 2D array `costArr`, where the rows i represent the # of edges in a path and the columns j represent the nodes in the graph. `costArr[i][j]` = cost of the shortest path from $v \rightarrow j$ of length i .
- Create a 2D array `numArr`, where the rows i represent the nodes and the columns j represent the # of edges. `numArr[i][j]` = number of shortest paths from $v \rightarrow i$ of length j .
- Set the `costArr[0][0...n] = 0` and `costArr[0...n][0] = 0`, for $v=0$ as the starting node.
- For # of edges $e = 1$ to $|E|$:
 - For nodes $x = 1$ to $|V|$:
 - Consider all neighboring nodes of x , say $\{x_1 \dots x_n\}$
 - Let $C = \text{costArr}[x_i][e-1] + \text{cost of the edge } (x, x_i)$.
 - Set `costArr[x][e]` = minimum value of C .
 - Let $S = \text{sum of numArr}[y_i][e-1]$, where $\{y_1 \dots y_n\}$ are all neighboring nodes that yield the optimal cost for path length $e-1$.
 - Set `numArr[x][e]` = S .
- To find the # of shortest paths from $v \rightarrow w$:
 - Iterate through `costArr[1...n][w]` and find the path lengths that give the shortest costs.
 - Sum up `numArr[w][p]` for each optimal path length $\{p_1 \dots p_n\}$
 - Return this sum.

Proof of Correctness

- Base case:
 - There is only 1 node in the graph, the starting node.
 - There are 0 shortest paths to the starting node.
 - Our algorithm establishes that the starting node will have 0 shortest paths to it by setting `costArr[0][0...n] = 0`.
- Inductive hypothesis:
 - Suppose we know the optimal (shortest) path length to all preceding vertices $\{v_1 \dots v_k\}$ of vertex w . That is, $(v_1 \dots v_k \rightarrow w)$ are valid directed edges in the graph.
 - We want to know the optimal path length to vertex w .
 - The optimal path length to w will be the shortest path length we can find by adding $[s \rightarrow \text{a preceding vertex } v]$ and $[\text{edge cost between } v \text{ and } w]$.
 - Because we have optimized the path length from $s \rightarrow v$, we know that it will be minimized.
 - Therefore, we just have to find the minimum of $(s \rightarrow v \text{ path length} + \text{edge cost})$, which our algorithm does.
 - By tracking the # of shortest paths to the vertex in a 2D array, we can effectively return the # of shortest path lengths to any vertex in the graph.

Time Complexity

- Iterating through all the edges in the graph in the outer loop takes $O(E)$ time.
- Iterating through each node in the inner loop takes $O(V)$ time.

- Therefore, the total runtime is $O(E \cdot V)$, where $E = \# \text{ edges}$ and $V = \# \text{ vertices}$.
- Space complexity is $O(V^2)$, because we are using 2D arrays.

3. Exercise 24 on page 331

Algorithm

- We can divide this problem into subproblems of gerrymandering precincts from 1-k, where $k = \{1 \dots n\}$. For every precinct we add, we can either place it in district 1 or district 2. We save the result of placing it in both district 1 and district 2, and continue onwards. This is because there's a possibility that the precinct is in either district for the optimal solution, so we must consider both cases.
- Create a 4D array $\text{arr}[i, j, k, l]$ with i representing the precinct #, j representing the # of precincts in district 1, k representing the # of A-votes in district 1, and l representing the # of A-votes in district 2.
- Initialize $\text{arr}[0, 0, 0, 0] = \text{true}$
- For $i = 1$ to n precincts:
 - For $j = 1$ to n precincts in district 1:
 - For $x = 0$ to $m \cdot n$ votes:
 - For $y = 0$ to $m \cdot n$ votes:
 - If either $\text{arr}[i-1, j-1, x-A[j], y]$ or $\text{arr}[i-1, j-1, x, y-A[j]]$ is true, then $\text{arr}[i, j, x, y] = \text{true}$. Otherwise, it's false.
- If there exists a true entry in arr where $\text{arr}[i, j, k, l] = \text{arr}[n, n/2, x, y]$ where x and y are both greater than $m \cdot n/4$, then we return true.
- Otherwise, return false.

Proof of Correctness

- Base case:
 - There are 2 precincts.
 - There will be a majority unless the sum of the 2 parties' votes are tied.
- Inductive hypothesis:
 - Suppose we know whether gerrymandering would work for the first 1-n precincts.
 - We want to know whether it would work for the n th precinct.
 - Given that we can either place this precinct in district 1 or 2, we can see which placement will yield a majority.
 - If placing it in district 1 or 2 yields a majority, then there would be a valid gerrymandering. If neither yields a majority, there is no valid gerrymandering.
 - Our algorithm keeps track of the results of these choices through the 4D array, and therefore is able to figure out whether gerrymandering works for the n th precinct.

Time Complexity

- We are dividing the problem down into $N^2 M^2$ subproblems, because there are N precincts and M total voters in each precinct.

- Because each computation in the innermost loop takes constant time, the total runtime of the algorithm is $O(N^2M^2)$.
- The space complexity is $O(N^4)$ because we are using a 4D array.

4. Exercise 7 on page 417

Algorithm

- Let n = # of clients and k = # of base stations.
- Create a network graph G with source node S , sink node T , and internal nodes that represent the clients and base stations.
- Add edges of capacity 1 between all client nodes v_c and the source node S .
- Add edges of capacity L (load parameter) between all station nodes v_s and the sink node T .
- Add an edge of capacity 1 between a client and station only if the client is within distance r of the station.
- If there exists a valid S - T flow with a value of n in G , then there is a way to connect all clients to a base station subject to the given requirements.

Proof of Correctness

- We've already proved the maxflow algorithm works.
- Consider a cut consisting of source S combined with all the edges it touches, and the rest of the nodes in the other partition.
- Because each internal edge has capacity 1, the number of edges (and therefore the number of clients) over this cut would equal the flow.
- Because each outgoing flow from the station nodes to the sink is L , each station node can be connected to at most L client nodes.
- Therefore, if there is an outgoing flow from the station nodes to the sink of L , that means that there is a total of L clients over the cut.
- That means that the max flow is n , and each client is matched to a station; our algorithm thus works to figure out whether all clients can be matched to a station.

Time Complexity

- The network graph G will have $n+k$ nodes and at most $n*k$ edges, so the runtime will be the time to solve such a max-flow problem.
- The max possible flow through G will be n , given that each client is connected to all stations with an edge of capacity 1.
- Finding a S - T path in the max-flow problem costs $O(V+E)$, or $O(n+k + n*k)$
- Therefore, the total runtime would be $O(n * (n+k + n*k))$, or roughly $O(n^2k)$.

5. Exercise 9 on page 419

Algorithm

- Let n = # of injured people and k = # of hospitals.
- Create a network graph G with source node S , sink node T , and internal nodes that represent the injured people and hospitals.
- Add edges of capacity 1 between all injured people nodes v_i and source node S .

- Add edges of capacity $\lceil n/k \rceil$ between hospital nodes v_h and sink node T .
- Add an edge of capacity 1 between an injured person v_i and hospital v_h only if the hospital is within half-hour's driving time of the person's current location.
- If there is a valid S-T flow with a value of n , then there exists a way to send all n people to a hospital without overloading any of the hospitals.

Proof of Correctness

- We've already proved the maxflow algorithm works.
- Consider a cut consisting of source S combined with all the edges it touches, and the rest of the nodes in the other partition.
- Because each internal edge has capacity 1, the number of edges (and therefore the number of injured people) over this cut would equal the flow.
- Because each outgoing flow from the hospital nodes to the sink is n , each hospital node can be connected to at most n people nodes.
- Therefore, if there is an outgoing flow from the hospital nodes to the sink of n , that means that there is a total of n injured people over the cut.
- That means that the max flow is n , and each person is matched to a hospital; our algorithm thus works to figure out whether everyone can be matched to hospitals.

Time Complexity

- The network graph G will have $n+k$ nodes and at most $n*k$ edges, so the runtime will be the time to solve such a max-flow problem.
- The max possible flow through G will be n , given that each person is connected to all hospitals with an edge of capacity 1.
- Finding a S-T path in the max-flow problem costs $O(V+E)$, or $O(n+k + n*k)$
- Therefore, the total runtime would be $O(n * (n+k + n*k))$, or roughly $O(n^2k)$.

6. **Given a sequence of numbers, find a subsequence of alternating order, find the length where the subsequence is as long as possible. (That is, find a longest subsequence with alternate low and high elements).**

Example

Input: 8, 9, 6, 4, 5, 7, 3, 2, 4

Output: 8, 9, 6, 7, 3, 4 (of length 6)

Explanation: 8 < 9 > 6 < 7 > 3 < 4 (alternating < and >)

Algorithm

- For this problem, we want to check the longest alternating subsequence for every possible ending element (index 0 to $n-1$). The ending element of the subsequence will either be a high (greater than the previous element) or a low (less than the previous element). We must consider both cases.
- Create a 2D array `arr` with rows 0- n that represent the ending index of the subsequence, and columns 0-1 that specify whether it's a low or high, respectively. `arr[i][0]` holds the length of the longest alternating subsequence that ends at index i and is a "low." `arr[i][1]` holds the length if the element is a "high."
- Set `arr[0][0]` and `arr[0][1]` to 1.

- Create a variable maxLen that tracks the maximum subsequence length.
- For each index i from 0 to n-1:
 - Keep highMax and lowMax variables that track the longest subsequence length if we end on a high or low, respectively.
 - For every preceding index j from 0 to i:
 - if $\text{arr}[j][1] + 1 > \text{lowMax}$ (case if we end on low)
 - set $\text{lowMax} = \text{arr}[j][1] + 1$
 - if $\text{arr}[j][0] + 1 < \text{highMax}$ (case if we end on low)
 - set $\text{highMax} = \text{arr}[j][0] + 1$
 - $\text{arr}[i][0] = \text{lowMax}$
 - $\text{arr}[i][1] = \text{highMax}$
 - Set $\text{maxLen} = \max(\text{lowMax}, \text{highMax})$ if it is greater than maxLen
- Return maxLen

Proof of Correctness

- The key to this algorithm is that the optimal alternating subsequence that ends on high = optimal preceding subsequence that ends on low + the current element. Likewise for the optimal subsequence that ends on low.
- Base case:
 - The alternating subsequence ends on index 0.
 - The max length out of all possible alternating subsequences = 1, which is just the element itself.
- Inductive hypothesis:
 - Suppose we know the length of the optimal alternating subsequences that end on or before indices n-1, for both whether the ending element is a low or high.
 - We want to figure out the length of the optimal alternating subsequence for the sequence 0-n (ends on index n).
 - If we have the max length of a preceding subsequence that ends on a low, then its length + 1 = length of optimal subsequence for sequence 0-n that ends on high.
 - If we have the max length of a preceding subsequence that ends on a high, then its length + 1 = length of optimal subsequence for sequence 0-n that ends on low.
 - The length of the optimal subsequence for sequence 0-n would therefore be the maximum out of these 2 values.
 - Our algorithm correctly finds the length of the alternating subsequence that ends on index n!

Time Complexity

- The outer loop goes through indices from 0-n, which takes $O(N)$ time. We are looping through every possible ending index.

- The inner loop goes through indices from 0-i, which takes worst case $O(N)$ time. We are looping through the preceding indices of i to find the most optimal length so far.
- Therefore, the total runtime is $O(N^2)$. Space complexity is $O(N)$ because we are using a 2D array that has a fixed column length of 2 but variable row length that depends on the input size.