

HW3 Rubric

1. Exercise 10 Page 110 (15 points)

In an undirected Graph $G=(V, E)$, given two nodes v and w , we want to find the number of all the shortest paths from v to w .

Algorithm (7 points)

- ☐ Run BFS from v with a slight modification. Store a counter for each node, initialized to 0, except for root node v , which is initialized to 1. We also need to store another variable in the nodes, denoting the layer they are located in, it is initialized to -1 except for the root (0), and gets updated when we discover nodes.
- ☐ Once we discovered a node y when exploring the neighbors of another node x , we do $\text{counter}[y] += \text{counter}[x]$, **only if y is in the next layer of x** .
 - ☐ If y was a node in the upper layer of x or its sibling, we shouldn't update $\text{counter}[y]$
- ☐ While running BFS, check if we have discovered the target node w . whenever a node at the layer of w was selected for exploration, we can terminate the algorithm.

Runtime analysis (3 points)

- ☐ Incrementing the counter and checking the layer can be done in $O(1)$. Hence, all we have is the normal BFS algorithm, which is of $O(m+n)$, plus some additional operations at each step, which could be done in $O(1)$, and don't change the overall runtime of our algorithm.

Proof (5 points)

- ☐ Use induction
 - ☐ Our algorithm correctly counts the number of paths for all the nodes in layers 1 to k .
- ☐ Base case: $k=1$: # of all the paths from v to v is 1
- ☐ Induction step: we have found all the paths from root (v) to layer $k-1$.
- ☐ For each node t in layer k :
 - ☐ In all the shortest paths from v to t , the node before t must be located in $(k-1)$ th layer. So, we have:

$$\text{Number of paths ending in } t = \sum_{\substack{z \in \text{layer } (k-1) \\ \text{and edge } (z,t) \text{ exists}}} \text{number of paths ending in } z$$

2. Exercise 6 on page 108 (15 points)

Proof: by contradiction: Assume there is an edge $e=(x, y)$ in G that does not exist in T . These cases might arise:

- x and y are more than 1 layers apart in T .
 - Contradiction: T is a BFS tree. When BFS is exploring one of the nodes (whichever it selects first) the other will be added to the next layer.
- x and y are exactly one layer apart
 - T is a DFS tree. without loss of generality, assume DFS selects x for exploration before y . x should be the ancestor of y , and because they are one layer apart, x is parent of y and e exists in T .
- x and y are siblings
 - Same as the previous case. Without loss of generality, assume x is explored first. y should be in the subtree of x , because T is a DFS tree. Can't be siblings. Contradiction.

All cases ran into a contradiction. Thus, $T=G$.


3. Exercise 12 on page 193 (20 points)

a) (5 points) Counterexample: suppose link parameter is $r=5000$, and we have $(b_1, t_1)=(1000,1)$ and $(b_2, t_2) = (6000, 1)$. Although $b_2 \leq rt_2$ does not hold here, running these two streams in the order of 1,2 is valid.

b) (15 points)

Claim: Given any n streams, the total number of transferred bits would be minimum at each time if we sort the video streams based on their bit rate $r_i = \frac{b_i}{t_i}$ and send them in non-decreasing order.


For better visualization, look at the following figure. The x and y axis are time and the total number of transferred bits respectively. the slope would be $m = \frac{b_2 - b_1}{t_2 - t_1}$ which has the same definition as bit transfer rate r . The question asks us to keep the diagram under the line with slope r .



Proof: by contradiction.

- assume the optimal ordering has an **inversion**. We define an inversion as two subsequent streams i, j for which $r_i > r_j$ (i is sent exactly before j). If the ordering is not non-decreasing, it is guaranteed that an inversion can be found.
 - Notation: the start time of sending i is t' and the time by which both i and j are sent is denoted by $t'' = t' + t_i + t_j$.

- Reverse the order of i and j . We claim that the resulting ordering would have a lower transfer rate at all times, except for the beginning and the end of the period.
- Look at the following figure. The lines marked with 1 (on the top) are regarding to when we send i before j . (2) is for the exact opposite. B is the total number of bits that have been sent by t' (before we start transferring i and j).



- Note we calculate overall bit rate by $r(t) = \frac{\text{\# of bits sent by } t}{t}$
 - For times $t < t'$ and $t > t''$, the bit transfer rate would be the same for both orderings.
 - for $t < t'$: Nothing has changed
 - for $t > t''$: We send $b_i + b_{i+1}$ bits in $[t', t'']$ in both cases: No difference in calculating $r(t)$
 - For $t \in [t', t'']$ the modified ordering has lower bit transfer rate, i.e., the total number of bits for each t in the mentioned period would be lower in path 2 in comparison with path 1 (in the diagram)
 - It can be easily shown by using the fact that $r_i > r_j$.

We improved the optimal algorithm and this is a contradiction. An inversion cannot exist and the optimal ordering must be non-decreasing.

In the first figure, it is clear that if the lines are sorted in a non-decreasing order, the whole figure would be under the line connecting the origin to the end point. Furthermore, we know that the slope of this line is $\frac{\sum_{k=1}^n b_i}{\sum_{l=1}^n t_l}$. Thus, to solve the problem it suffices to calculate this fraction, which can be done in $O(n)$. Note that giving the exact order of sending video streams still needs $O(n \log(n))$ time, but it is not requested by the question.

Detailed rubric for part b: definition of algorithm: 5pts, proof: 5pts, Runtime: 5pts

4. Exercise 3 on page 189 (15 points)

We want to prove the greedy algorithm always stays ahead. Same as what we did in the interval scheduling problem, we use proof by contradiction: assume greedy algorithm is not optimal and another algorithm is doing better than it. Use induction for comparison:

- ☐ The greedy algorithm sends more or the same number of boxes in the first k trucks.
- ☐ Base case is obvious.
- ☐ Suppose the statement is correct for $k-1$.
 - Greedy algorithm have sent b_1, \dots, b_i , and the other one b_1, \dots, b_j , where $i \geq j$
- ☐ For k^{th} round:
 - The other algorithm sends b_{j+1}, \dots, b_l . We know that the sum of the weights of b_{i+1}, \dots, b_l is less than W because $i \geq j$. So we can first put the packages that the other algorithm has sent in the k^{th} truck, and load the remaining space with other packages.
 - So, our algorithm is ahead after k rounds.

Overall, the greedy approach uses less number of trucks than other algorithms, or at worst case, equal.

5. Exercise 6 on page 191 (20 points)


Algorithm (7 points)

Suppose s_i, b_i, r_i are projected swimming, biking, and running time for contestant i . We sort all contestants based on $b_i + r_i$ in non-increasing order and claim this is the optimal order.


Proof (10 points)

Imagine the optimal order is not what we suggested. The optimal algorithm does not follow the non-increasing order in terms of $b + r$, so we can find two subsequent contestants i, j (j starts immediately after i) for which $b_j + r_j > b_i + r_i$. We claim if we swap their position, the resulting order would be better.

- ☐ The pool is in use for $s_i + s_j$ in both orders.
- ☐ The finishing time of i and j in the modified order w.r.t. the time they started using the pool is:
 - $\max(s_j + b_j + r_j, s_j + s_i + b_i + r_i)$



- ☐ in the original order, it was:
 - $s_i + s_j + b_j + r_j$



Because $b_i + r_i > b_j + r_j$, the finishing time of the modified algorithm is better, and this is a contradiction to the original assumption: it must be impossible to improve the optimal algorithm.

Hence, an algorithm with an inversion cannot be optimal, and our claim was correct.

Time complexity (3 points)

Calculating $r+b$ for all contestants is $O(n)$. Sorting could be done in $O(n \log(n))$. Overall, $O(n \log(n))$.

6. The rotting orange problem (15 points)

{2, 1, 0, 2, 1}

{1, 0, 1, 2, 1}

{1, 0, 0, 2, 1}

Algorithm (7 points)

The solution would be similar to running several BFS algorithms simultaneously.

In the beginning, put the indexes of all the 2s in a queue and call it $q[0]$.

- ☐ $q[0] = [(0, 0), (0, 3), (1, 3), (2, 3)]$.
- ☐ $q[i]$ is the set of all oranges that rot in timestamp i

Moreover, count the total number of ones in the matrix and call it $cnt1$.

Follow this algorithm:

- ☐ Pop the first element of $q[i]$
- ☐ Add the indexes of all the 1s adjacent to it to $q[i+1]$
- ☐ Change those 1s to 2
- ☐ Subtract the number of added elements from $cnt1$
- ☐ Repeat the following steps until $q[i]$ is empty, then If $q[i+1]$ is empty
 - ☐ If $cnt1=0$, it means that all the oranges are rotten at the end. Return i
 - ☐ If $cnt1>0$: it is impossible for every orange to be rotten. Return -1

Time complexity analysis (5 points)

If we take each index as a node, the number of nodes is NM . We select each node for exploration at most once and explore 4 directions. The overall runtime would be $O(MN)$.

Proof (3 points)

We should prove if there exists a shortest path p with length m from a rotten orange to a fresh orange:

- ☐ The fresh orange will definitely rot
- ☐ It will rot at time m

Both are obvious based on how BFS works. Moreover, it is obvious that if there is no path to a fresh orange, the algorithm will return -1.